

Session ID	Day	Session Start Time	Room	Session Title	Session Capacity	Paper Order In Session	Presentation Start	Presentation End	Paper ID	Cycle	Paper Title	Authors	Notes
S01	Tuesday	09:20 AM	Room A	Best Paper Candidates	3	1	9:20 AM	9:40 AM	184	2	Coordinating GPU Data Centers and Power Grid Regulation Service for Exogenous Carbon Benefits	D. Wong, A. Jahanshahi, S. Golrouye, O. Anderson, N. Yu	
S01	Tuesday	09:20 AM	Room A	Best Paper Candidates	3	2	9:40 AM	10:00 AM	211	1	FLYING SERVING: On-the-Fly Parallelism Switching for Large Language Model Serving	S. Gao, J. Yin, F. Wang, W. Dong	
S01	Tuesday	09:20 AM	Room A	Best Paper Candidates	3	3	10:00 AM	10:20 AM	785	2	OCTANE: Breaking the Neighbor-List Bottleneck in GPU Molecular Dynamics	H. Toutouni, S. Chakraborty, Y. Tu, J. Huang	
S02	Tuesday	10:50 AM	Room A	Compiler, Code Generation and Autotuning	4	1	10:50 AM	11:10 AM	74	2	GRASP: Optimizing VLIW Instruction Scheduling via Graph Reinforcement	Z. Wang, W. Tong, J. Fang, Y. Zhang, W.	
S02	Tuesday	10:50 AM	Room A	Compiler, Code Generation and Autotuning	4	2	11:10 AM	11:30 AM	283	2	Continuation-Preserving Tiling for Pointer-Chasing Optimization in Structured	A. Kumar, V. Singh, S. Biswas	
S02	Tuesday	10:50 AM	Room A	Compiler, Code Generation and Autotuning	4	3	11:30 AM	11:50 AM	427	2	S2VEC: Compiler-Driven Stream Specialization for Linearized Vectorization	L. Crespo, A. Fernandes, G. Falcao, P.	
S02	Tuesday	10:50 AM	Room A	Compiler, Code Generation and Autotuning	4	4	11:50 AM	12:10 PM	557	1	CKT1: A Domain-Specific Compiler for Lowering CUDA Kernels to Triton-IR	C. Shi, R. Chen, Y. Sun, Y. Sui, J. Zhang,	
S03	Tuesday	10:50 AM	Room B	Runtime Scheduling and Adaptive Execution	4	1	10:50 AM	11:10 AM	66	2	Barrier-Aware Task Scheduling for Bulk-Synchronous Parallel Architectures	T. Noack, A. Koch	
S03	Tuesday	10:50 AM	Room B	Runtime Scheduling and Adaptive Execution	4	2	11:10 AM	11:30 AM	435	1	FaaSIm: Partial Caching of Snapshot-based VMs for Serverless Computing	S. Eom, C. Park, G. Lee, H. Moon, Y. Choi	
S03	Tuesday	10:50 AM	Room B	Runtime Scheduling and Adaptive Execution	4	3	11:30 AM	11:50 AM	195	2	Block-Aware Adaptive State Management for Optimistic Parallel Discrete Event Simulation	X. Peng, Q. Wang, G. Liu, C. Hong, R. Xia, Z. Sun, X. Chen, Q. Zhang, J. Liu	
S03	Tuesday	10:50 AM	Room B	Runtime Scheduling and Adaptive Execution	4	4	11:50 AM	12:10 PM	422	1	Lock Shielding: A General Technique for Misuse-Resilient Locks	V. Shahare, M. Chabbi, N. Hegde	
S04	Tuesday	10:50 AM	Room C	Energy & Sustainability	4	1	10:50 AM	11:10 AM	63	1	Watchmen: Watching the Watchers – High Fidelity, Flexible GPU Energy Modeling	B. Tran, M. Sinclair, S. Venkataraman, M. Matherly, W. Shin	
S04	Tuesday	10:50 AM	Room C	Energy & Sustainability	4	2	11:10 AM	11:30 AM	89	1	Agile QoS-aware Dynamic Power Management with eBPF Governors	M. Rezvani, D. Wong	
S04	Tuesday	10:50 AM	Room C	Energy & Sustainability	4	3	11:30 AM	11:50 AM	212	1	SmartCap: Coordinated CPU-GPU Power Capping for Performance-Assurance Energy Efficiency	Z. Zheng, Z. Lan, X. Wu, V. Taylor, M. Papka	
S04	Tuesday	10:50 AM	Room C	Energy & Sustainability	4	4	11:50 AM	12:10 PM	464	1	CATS: Correlation-aware Task Scheduling for GPU Power Optimization in AI Data Centers	S. Subramanian, X. Wang	
S05	Tuesday	01:40 PM	Room B	Performance Modeling & Insight	4	1	1:40 PM	2:00 PM	74	1	TenProf: A Tensor-Centric Profiler for Deep Learning Workload Analysis and Optimization	X. Ding, K. Zhou, Y. Hao, P. Su	
S05	Tuesday	01:40 PM	Room B	Performance Modeling & Insight	4	2	2:00 PM	2:20 PM	82	1	GRASP: Fine-grained and Adaptive Sampled Simulation for GPU Performance Modeling	L. Chao, Z. Huang, P. Cai, J. Xue, T. Xiong, R. Xue	
S05	Tuesday	01:40 PM	Room B	Performance Modeling & Insight	4	3	2:20 PM	2:40 PM	217	1	Mantis: Decoding HPC Telemetry Data for Robust System Prediction	Y. Lu, J. Ren, S. Evgenia	
S05	Tuesday	01:40 PM	Room B	Performance Modeling & Insight	4	4	2:40 PM	3:00 PM	594	2	VSim: A Lightweight SpMV Performance Simulator via Statistical and Visual Residual Learning	S. Zhu, W. Huangfu, G. Chu	
S06	Tuesday	01:40 PM	Room C	I/O & Storage	4	1	1:40 PM	2:00 PM	12	1	Harmonia: Enhancing Data Placement and Migration in Hybrid Storage Systems via Multi-Agent Reinforcement Learning	R. Nadig, V. Arulchelvan, R. Bera, T. Shahroodi, G. Singh, A. Kakolyris, I. Yuksel,	
S06	Tuesday	01:40 PM	Room C	I/O & Storage	4	2	2:00 PM	2:20 PM	177	2	ColdMap: Compaction-Aware Cost-Benefit Zone Cleaning for ZNS-Based Key-	S. Byeon, K. Min, J. Park, S. Lee, H. Kim, J.	
S06	Tuesday	01:40 PM	Room C	I/O & Storage	4	3	2:20 PM	2:40 PM	285	1	CoCache: Accelerating Reads in KV Stores via Cooperative Metadata and Data	H. Tang, W. Zhu, Q. Zhang, J. Zhang, J.	
S06	Tuesday	01:40 PM	Room C	I/O & Storage	4	4	2:40 PM	3:00 PM	446	2	TOTO: Transparent I/O Tuning for HPC Applications	F. Boito, L. Teylo, M. Popov, L. Aimi, A.	
S07	Wednesday	10:50 AM	Room A	Graph Search & Paths	4	1	10:50 AM	11:10 AM	44	1	G-PathGen: An Efficient GPU-Parallel k-Critical Path Generation Algorithm	C. Chang, Y. Chung, C. Chiu, W. Lee, B.	
S07	Wednesday	10:50 AM	Room A	Graph Search & Paths	4	2	11:10 AM	11:30 AM	216	1	Parallel Bidirectional A* Search for GPU-Accelerated Pathfinding	H. Al Khansa, J. Luna, A. Mouawad, I. Hajj	
S07	Wednesday	10:50 AM	Room A	Graph Search & Paths	4	3	11:30 AM	11:50 AM	469	1	MPMOS: Massively Parallel Multi-Objective Shortest Paths	L. Gold, D. Sidoti, K. Pattipati, O. Khan	
S07	Wednesday	10:50 AM	Room A	Graph Search & Paths	4	4	11:50 AM	12:10 PM	844	1	Distributed Disjoint Weighted Matchings in Demand-Aware Reconfigurable Optical	S. Heck, K. Hanauer, S. Schmid	
S08	Wednesday	10:50 AM	Room B	AI Training Systems	4	1	10:50 AM	11:10 AM	68	2	Closing the Efficiency Gap: AI Datacenter Co-design Roadmap for Scalable	J. Tihi, H. Wu, J. Park, A. Abuhatzera, F.	
S08	Wednesday	10:50 AM	Room B	AI Training Systems	4	2	11:10 AM	11:30 AM	131	2	COMETS: Cost-effective Multi-node Efficient Training System with Memory	H. Chen, S. Yang, M. Soltaniyeh, S. Pei, A.	
S08	Wednesday	10:50 AM	Room B	AI Training Systems	4	3	11:30 AM	11:50 AM	282	1	SPPO: Making Million-Token LLM Training Practical on Modest GPU Clusters	q. chen, S. Li, W. GAO, P. Sun, Y. Wen, T.	
S08	Wednesday	10:50 AM	Room B	AI Training Systems	4	4	11:50 AM	12:10 PM	392	1	Rudder: Steering Prefetching in Distributed GNN Training using LLM Agents	A. Sarkar, S. Ghosh, N. Tallent, A. Chadha,	
S09	Wednesday	10:50 AM	Room C	Communication & Collectives	4	1	10:50 AM	11:10 AM	19	2	SHIRO: Near-Optimal Communication Strategies for Distributed Sparse Matrix	C. Zhuang, L. Zhang, B. Brock, D. Wu, P.	
S09	Wednesday	10:50 AM	Room C	Communication & Collectives	4	2	11:10 AM	11:30 AM	187	2	PACER: A Userspace Network Rate Controller in MPI with Adaptive Compression	Y. Li, D. Ng, A. Kashyap, S. Di, G. Li, X. Lu	
S09	Wednesday	10:50 AM	Room C	Communication & Collectives	4	3	11:30 AM	11:50 AM	387	1	Scalable All-to-all Algorithms for Dynamic and Irregular Communication Patterns	C. Wei, A. Bhatle	
S09	Wednesday	10:50 AM	Room C	Communication & Collectives	4	4	11:50 AM	12:10 PM	78	1	StencilMD: Optimizing Communication in Molecular Dynamics Simulations	R. Deng, T. Schardl	
S10	Wednesday	01:40 PM	Room A	Graph Analytics	4	1	1:40 PM	2:00 PM	47	1	DCSM: Enabling Inter-Batch Parallelism for Continuous Subgraph Matching on	Y. Wei, P. Jiang	
S10	Wednesday	01:40 PM	Room A	Graph Analytics	4	2	2:00 PM	2:20 PM	160	2	SVSIG: Incremental Streaming Graph Processing with Source Vertex Suppression	J. Huang, X. Yan, D. Fu, H. Bian, T. Cao, Z.	
S10	Wednesday	01:40 PM	Room A	Graph Analytics	4	3	2:20 PM	2:40 PM	193	2	GAAF: Fast and Scalable Graph-based Vector Similarity Search with Any-Match	M. Ma, X. Yin, J. Qiu	
S10	Wednesday	01:40 PM	Room A	Graph Analytics	4	4	2:40 PM	3:00 PM	344	2	HoloGraph: Bridging the Throughput Gap in Heterogeneous Graph Pattern	M. Haotian, W. Hsu, Y. Chung	
S11	Wednesday	01:40 PM	Room B	LLM Serving	4	1	1:40 PM	2:00 PM	155	1	Taming Dynamic Diffusion LLM Inference through Virtual Static Execution	j. Zhu, H. Wu, Y. Li, H. Wang, R. Li, J. ZHAI	
S11	Wednesday	01:40 PM	Room B	LLM Serving	4	2	2:00 PM	2:20 PM	222	1	SparseServe: Unlocking Parallelism for Dynamic Sparse Attention in Long-Context	Q. Zhou, P. Yin, P. Zuo, C. Wang, J. Cheng	
S11	Wednesday	01:40 PM	Room B	LLM Serving	4	3	2:20 PM	2:40 PM	231	1	InferFast: Bridging the Gap Between Unstructured LLM Sparsity and Practical GPU	Z. Shen, W. Bu, X. He, K. Sheng, H. Chen	
S11	Wednesday	01:40 PM	Room B	LLM Serving	4	4	2:40 PM	3:00 PM	268	1	dLLM-Serve: Bridging the Memory Gap in Diffusion Language Model Serving	J. Fan, Y. Zhang, X. Li, D. Nikolopoulos	
S12	Wednesday	01:40 PM	Room C	Near-Memory Architectures	4	1	1:40 PM	2:00 PM	38	1	DANMP: Accelerating Multi-Scale Deformable Attention Using Near-Memory-	H. Li, Q. Wang, B. Gao, D. Chen, Y. Huang,	
S12	Wednesday	01:40 PM	Room C	Near-Memory Architectures	4	2	2:00 PM	2:20 PM	190	2	RPFC: A Router Partitioning and Forward Channel Routing Framework for 2.5D	S. Tao, Z. Guo, T. Liu, J. Wang	
S12	Wednesday	01:40 PM	Room C	Near-Memory Architectures	4	3	2:20 PM	2:40 PM	272	1	MyT: Efficient Manycore based on Many Threading and Scalable Memory	A. Rajasukumara, R. Xu, T. Zhang, Y. Wang,	
S12	Wednesday	01:40 PM	Room C	Near-Memory Architectures	4	4	2:40 PM	3:00 PM	521	1	Clutch: High Performance Vector-Scalar Comparison using DRAM via Chunked	D. Tokuda, T. Kubo, I. Yuksel, A. Olgun, H.	
S13	Wednesday	03:30 PM	Room A	Graph Traversal & Connectivity	5	1	3:30 PM	3:50 PM	225	1	cuMIS: A Unified Scalable Framework for Computing Maximal Independent Sets	J. Nke, S. Kang, B. Rees, C. Lee	
S13	Wednesday	03:30 PM	Room A	Graph Traversal & Connectivity	5	2	3:50 PM	4:10 PM	243	1	BLEST: Blazingly Efficient BFS using Tensor Cores	D. Elbek, K. Kaya	
S13	Wednesday	03:30 PM	Room A	Graph Traversal & Connectivity	5	3	4:10 PM	4:30 PM	485	1	CORE-BFS: Communication Optimized Rectangular Partitioned BFS Achieving	H. Yang, H. Lu, M. Matheson, F. Wang, H.	
S13	Wednesday	03:30 PM	Room A	Graph Traversal & Connectivity	5	4	4:50 PM	5:10 PM	668	2	DynLP: Parallel Dynamic Batch Update for Label Propagation in Semi-Supervised	S. Shovan, A. Khanda, S. Ferdous, S. Das,	
S13	Wednesday	03:30 PM	Room A	Graph Traversal & Connectivity	5	5	4:30 PM	4:50 PM	604	2 (sheph)	Asynchronous Distributed Algorithms for Maximal-Weighted Graph Matching	S. Singhal, J. Yang, A. Hayashi, V. Sarkar	
S14	Wednesday	03:30 PM	Room B	Data Analytics & Compression	5	1	3:30 PM	3:50 PM	186	1	GPZ: GPU-Accelerated Lossy Compressor for Particle Data	R. Li, Y. Huang, L. Zhang, Z. Yang, S. Di, B.	
S14	Wednesday	03:30 PM	Room B	Data Analytics & Compression	5	2	3:50 PM	4:10 PM	593	2	GFAz: State-of-the-Art Graphical Fragment Assembly Compression	T. Yang, Y. Liu, B. Jiang, X. Shi, S. Jin	
S14	Wednesday	03:30 PM	Room B	Data Analytics & Compression	5	3	4:10 PM	4:30 PM	425	2	Optimizing Streaming Tensor Decomposition on GPU	W. Lin, J. Sheng, S. Feng, M. Dun, H. Cao,	

S14	Wednesday	03:30 PM	Room B	Data Analytics & Compression	5	4	4:30 PM	4:50 PM	639	2	DA-MLAD: Drift-Decomposed Meta-Learning for Continual Log Anomaly Detection	K. Tan, Y. Du, D. Zhan, Y. Xie, H. Yu, B.
S14	Wednesday	03:30 PM	Room B	Data Analytics & Compression	5	5	4:50 PM	5:10 PM	32	2	TADS: Trend-Aware Dynamic Load Balancing for Large-Scale SNN Simulations	H. Huang, S. Pang, Y. Zeng, G. Feng, Z.
S15	Wednesday	03:30 PM	Room C	CXL & Memory Systems	5	1	3:30 PM	3:50 PM	134	2	CXL-CCL: Inter-Node Collective GPU-Communication Using a CXL Shared	D. Xu, H. Meng, X. Chen, D. Zhu, W. Tang,
S15	Wednesday	03:30 PM	Room C	CXL & Memory Systems	5	2	3:50 PM	4:10 PM	159	1	IBEX: Internal Bandwidth-Efficient Compression Architecture for Scalable CXL	Y. Ko, H. Park, H. Lee, H. Lee
S15	Wednesday	03:30 PM	Room C	CXL & Memory Systems	5	3	4:10 PM	4:30 PM	163	1	Clone: A Collaborative Multi-device System for Retrieval-Augmented Generation	S. Ko, W. Doh, E. Na, H. Shim, S. Yun, J.
S15	Wednesday	03:30 PM	Room C	CXL & Memory Systems	5	4	4:30 PM	4:50 PM	171	1	Anchoring Whole-System Persistence and Resilience in CXL	Y. Zhou, J. Zeng, C. Jung
S15	Wednesday	03:30 PM	Room C	CXL & Memory Systems	5	5	4:50 PM	5:10 PM	784	1	Griffin: Coherency-Aware Task Scheduling and Memory Allocation for CXL	S. Lee, K. Diab, D. Tootaghaj, L. Cao, P.
S16	Thursday	09:20 AM	Room A	Cross-Layer Performance Optimization	3	1	9:20 AM	9:40 AM	151	2	THAC: Unlocking Performance in Parallel HPC Applications via UQ-Aware	z. zhao, b. wang, B. yang, X. Chen, J. Liu,
S16	Thursday	09:20 AM	Room A	Cross-Layer Performance Optimization	3	2	9:40 AM	10:00 AM	147	1	Cross-Architecture Autotuning for Single-Source Heterogeneous Programming	H. Abram, N. Papadopoulou, J. Domke, M.
S16	Thursday	09:20 AM	Room A	Cross-Layer Performance Optimization	3	3	10:00 AM	10:20 AM	90	2 (sheph)	Look Before You Leap : Precision Instruction Supply via SmartScout	X. Zhang, P. Qu, T. Zhang, F. Su, Z. Pan,
S17	Thursday	10:50 AM	Room A	GPU-Accelerated Query and Geometry	4	1	10:50 AM	11:10 AM	555	2 (sheph)	Parallel Query Processing through Optimal Key Grouping on GPU-Based B+-Trees	Z. Chen, J. Li, J. Meng, N. Pitaksiranan, Y. Tu, B. Zeng, C. Dong
S17	Thursday	10:50 AM	Room A	GPU-Accelerated Query and Geometry	4	2	11:10 AM	11:30 AM	27	1	X-HD: Fast Hausdorff Distance Computation with Ray Tracing	L. Geng, Z. Yuan, R. Lee, X. Zhang, F. Wang
S17	Thursday	10:50 AM	Room A	GPU-Accelerated Query and Geometry	4	3	11:30 AM	11:50 AM	53	2	Rethinking Collision Detection on GPU Ray Tracing Architecture	D. Mandarapu, I. Fuksman, A. Pelenitysn, G. Bernstein, M. Kulkarni
S17	Thursday	10:50 AM	Room A	GPU-Accelerated Query and Geometry	4	4	11:50 AM	12:10 PM			EMPTY SLOT	Intentional empty slot
S18	Thursday	10:50 AM	Room B	AI Inference	4	1	10:50 AM	11:10 AM	23	2	LayerScope: Predictive Cross-Layer Scheduling for Efficient Multi-Batch MoE	E. Yu, D. Dong, Z. Zhang, Z. Bai, W. Yang,
S18	Thursday	10:50 AM	Room B	AI Inference	4	2	11:10 AM	11:30 AM	281	1	Aurora: A Disaggregated GPU-PNM-PIM System for High-Throughput Mixed-	H. Kim, S. Yu, M. Kim, J. Lee, H. Sung, E.
S18	Thursday	10:50 AM	Room B	AI Inference	4	3	11:30 AM	11:50 AM	292	1	HOPO: Accelerating Multimodal Neural Networks Inference via Holistic Parallelism	Y. Zheng, J. Sun, H. Li, G. Sun, J. Li
S18	Thursday	10:50 AM	Room B	AI Inference	4	4	11:50 AM	12:10 PM	650	2	Latency-SLO-Aware Memory Offloading for Large Language Model Inference	C. Ma, H. Zhao, Z. Ye, Z. Yang, T. Fu, J.
S19	Thursday	10:50 AM	Room C	Resilience and Error Detection	4	1	10:50 AM	11:10 AM	56	1	Not All Errors Are Equal: A Systematic Study of Error Propagation in Large	Y. Huang, S. Di, G. Li
S19	Thursday	10:50 AM	Room C	Resilience and Error Detection	4	2	11:10 AM	11:30 AM	366	2	SpinTune: Improving the Reliability of Quantum Sensor Networks for Practical	J. Ludmir, N. DiBrita, J. Han, P. Tirthak
S19	Thursday	10:50 AM	Room C	Resilience and Error Detection	4	3	11:30 AM	11:50 AM	388	2	Hamessing MPI mutations for AI error detection	A. Auville, T. Jammer, E. Pettit, P. Castro,
S19	Thursday	10:50 AM	Room C	Resilience and Error Detection	4	4	11:50 AM	12:10 PM	610	2	StreamGuard: Low-Overhead Resilience for Real-time HPC Data Streams	H. Nguyen, B. Nicolae, T. Bicer, A.
S20	Thursday	01:40 PM	Room A	Sparse & Tensor Kernels	4	1	1:40 PM	2:00 PM	25	1	Parametric Mappings for Distributed Tensor Computations	B. Wu, M. Kong
S20	Thursday	01:40 PM	Room A	Sparse & Tensor Kernels	4	2	2:00 PM	2:20 PM	415	1	Communication-Avoiding SpGEMM via Trident Partitioning on Hierarchical GPU	J. Bellavita, L. Pichetti, T. Pasquali, F.
S20	Thursday	01:40 PM	Room A	Sparse & Tensor Kernels	4	3	2:20 PM	2:40 PM	501	2	Ocean: Fast Estimation-Based Sparse General Matrix-Matrix Multiplication on GPU	Y. Li, G. Guidi
S20	Thursday	01:40 PM	Room A	Sparse & Tensor Kernels	4	4	2:40 PM	3:00 PM	277	1	PolyKAN: A High-Performance and Universal GPU Operator Library for Polynomial	m. yu, H. Zhong, J. Jiang, D. Huang, Y. Lu
S21	Thursday	01:40 PM	Room B	AI Kernels & Parallelism	4	1	1:40 PM	2:00 PM	269	1	HPMD: Enabling Hybrid Parallelism with Multi-Dimensional Adaptive DNN Training	G. Yun, Y. Choi
S21	Thursday	01:40 PM	Room B	AI Kernels & Parallelism	4	2	2:00 PM	2:20 PM	381	1	EPLoN: Exploiting Efficient Parallelism with Selective Rematerialization for	H. Bao, Z. Su, A. Setyae, S. Kamenev, A.
S21	Thursday	01:40 PM	Room B	AI Kernels & Parallelism	4	3	2:20 PM	2:40 PM	475	1	DynSpAttn: Efficient Attention via Dual-Side Dynamic Sparsity on Sparse Tensor	R. FAN, X. YU, Z. Li, W. Luo, G. Gong, X.
S21	Thursday	01:40 PM	Room B	AI Kernels & Parallelism	4	4	2:40 PM	3:00 PM	766	2	Three Birds, One Stone: Fast, Accurate-aware and Cost-Efficient Accelerator for	W. Jung, J. Kang, S. Shin, H. Um, J. Lim,
S22	Thursday	01:40 PM	Room C	Energy-Aware Systems	4	1	1:40 PM	2:00 PM	144	2	DEFT: Joint Task Placement and DVFS for Energy-Efficient Multi-GPU Runtimes	J. Chen, M. Pericás
S22	Thursday	01:40 PM	Room C	Energy-Aware Systems	4	2	2:00 PM	2:20 PM	278	2 (sheph)	Phase-aware Peak Power Reduction for Minimizing the Capital Expense of LLM	S. Wu, Y. Ma, X. Wang
S22	Thursday	01:40 PM	Room C	Energy-Aware Systems	4	3	2:20 PM	2:40 PM	584	1	Exploiting Hybrid Energy Storage to Minimize the Carbon Footprint of AI Data	S. Wu, X. Wang
S22	Thursday	01:40 PM	Room C	Energy-Aware Systems	4	4	2:40 PM	3:00 PM	307	2	The Performance-Power Frontier: A Model-Driven Approach to Energy-Aware	S. Pasupuleti, S. Wright
S23	Thursday	03:30 PM	Room A	Numerical & Scientific Kernels	5	1	3:30 PM	3:50 PM	159	2	WindStencil: Unleashing GPU Potential for High-Order Stencil Computation in High-	X. Zhang, H. Zhang, X. Liu, J. Li, R. Jin, J.
S23	Thursday	03:30 PM	Room A	Numerical & Scientific Kernels	5	2	3:50 PM	4:10 PM	252	2	Cheetah: Optimizing Execution Pipelines for Matrix-Free Finite Element Operators	J. Ren, H. Ltaief, S. Zampini, D. Keyes
S23	Thursday	03:30 PM	Room A	Numerical & Scientific Kernels	5	3	4:10 PM	4:30 PM	612	1	AdaPolySI: Adaptive Polynomial Filtered Subspace Iteration for Hermitian Interior	Y. Ni, X. Xu, S. Li, J. Zhang, J. Chen, J.
S23	Thursday	03:30 PM	Room A	Numerical & Scientific Kernels	5	4	4:30 PM	4:50 PM	679	1	Non-Delayed Cholesky Factorization	Y. Luo, S. Zhang, W. Liu
S23	Thursday	03:30 PM	Room A	Numerical & Scientific Kernels	5	5	4:50 PM	5:10 PM	87	2	Parallel Quadratic Selected Inversion in Quantum Transport Simulation	V. Maillou, M. Bollhofer, O. Schenk, A.
S24	Thursday	03:30 PM	Room B	Efficient Privacy Computing	5	1	3:30 PM	3:50 PM	95	2	MegaZK: A Memory Efficient GPU System Accelerating End-to-end Zero-	M. Li, Y. Yu, B. Wang, X. Fan, M. Gao, S.
S24	Thursday	03:30 PM	Room B	Efficient Privacy Computing	5	2	3:50 PM	4:10 PM	107	2	SumcheckPIM: An Efficient HBM-Based PIM Architecture for Linear Complexity	C. Kim, T. Kang, S. Shin, T. Suh, Y. Yang,
S24	Thursday	03:30 PM	Room B	Efficient Privacy Computing	5	3	4:10 PM	4:30 PM	117	2	GPiR: Enabling Practical Private Information Retrieval with GPUs	H. Ji, H. Yu, J. Kim, W. Choi, G. Suh, J.
S24	Thursday	03:30 PM	Room B	Efficient Privacy Computing	5	4	4:30 PM	4:50 PM	302	1	Scaling Long-Sequence Homomorphic Encrypted Transformer Inference via Hybrid	Z. Gong, R. Ran, F. Yao, W. Wen
S24	Thursday	03:30 PM	Room B	Efficient Privacy Computing	5	5	4:50 PM	5:10 PM	685	2	CipherSkip: Efficient Sparse Matrix Multiplication with FHE	W. Xiong, H. Zhou, Y. Ye, R. Jin, L. Xu
S25	Thursday	03:30 PM	Room C	Quantum Computing	5	1	3:30 PM	3:50 PM	48	2 (sheph)	qeStab: Towards Scalable Quantum Circuit Simulation on Multi-GPU using an	H. Shin, S. Lee, Y. Kim
S25	Thursday	03:30 PM	Room C	Quantum Computing	5	2	3:50 PM	4:10 PM	67	2	EZCache: A Hierarchical Memory System for Zoned Neutral Atom Quantum	J. Zhong, Y. Deng, H. Jiang, J. Feng
S25	Thursday	03:30 PM	Room C	Quantum Computing	5	3	4:10 PM	4:30 PM	375	2	TuniQ: Autotuning Compilation Passes for Quantum Workloads at Scale for	M. Hasanat, J. Ludmir, T. Patel, R. Roy
S25	Thursday	03:30 PM	Room C	Quantum Computing	5	4	4:30 PM	4:50 PM	424	2	C-3PQ: A Closeness Centrality-based Circuit Partitioner for Quantum Simulations	D. Popovici, H. Lee, N. Yoshioka, M. Ben,
S25	Thursday	03:30 PM	Room C	Quantum Computing	5	5	4:50 PM	5:10 PM	524	2	Diagonal-Budgeted Trotterization for Efficient Quantum Hamiltonian Simulation	S. Chundury, B. Burgstahler, J. Li, I. Suh, F.